

Runtime Process Infection

Shawn Webb, shawn@shawnwebb.info

September 19, 2012

Abstract

Injecting arbitrary executable code into a process is the end goal of exploitation. Many techniques already exist for injecting code. This paper will describe existing techniques as well as a novel new technique. The new technique provides stealthy injection into anonymous memory maps. The reader will need to be well-versed in ELF, Linux memory management on 32bit Intel systems, and the C programming language. Basic knowledge of assembly on 32bit Intel systems is assumed. Though the paper describes 32bit systems, a reference tool has been implemented that works on both 32bit and 64bit Linux systems.

1 Introduction

Writing malware on Windows is a relatively easy task. The win32 API is very friendly towards malware authors. Microsoft even provides a library called detours which makes injecting DLLs into a process during runtime extremely easy. Detours, however, does not provide stealth. Linux has a similar project called injectso. Injectso is a non-stealthy solution for injecting shared objects into a process during runtime.

If a system administrator thinks a service running on a server has been exploited, one thing the administrator can do is look at the `/proc/[pid]/maps` file and look at what shared objects are loaded. As an example, let's say we're targeting the Apache web server. We want to drop into a shell whenever the string `"GET /shell HTTP/1.1"` is sent to Apache over the socket. If we were to use injectso, we'd inject a shared object in such a way that the administrator

would see an entry that matches `/path/to/evil.so`. Obviously, the system's administrator knows at that point the game is over. And so will you when you try to get back in.

Therefore, we need a stealthy way to inject our shared object. Our requirements are:

1. Inject shared objects
2. Hook dynamically-loaded functions
3. Stealth through anonymous injection

In the rest of this paper, we will discuss:

1. Current techniques for storing and executing arbitrary shellcode
2. The Executable and Linkable Format (ELF)
3. The new technique for storing and executing arbitrary shellcode
4. Hooking dynamically-loaded functions
5. Future work and research

2 Current Techniques

Many techniques exist for storing and executing arbitrary code. For decades, storing code on the stack has been a known technique¹. Being well known, this technique is very well protected against². We can store our shellcode on the heap. Using the heap would be really simple since we can easily expand it

¹See aleph1's famous Phrack article, Smashing The Stack For Fun And Profit

²Add footnote for non-exec stack

with the `brk` system call. This technique is relatively old as well. Many systems are setting the heap to be non-executable. With the stack and the heap being protected, the other option would be to store our code where `$eip` points to. This would be an ideal solution, except that we need to hook dynamically-loaded functions. Meaning, by overwriting the code at `$eip`, we destroy any chance of the program using that spot in memory again unless we backup and restore the original code. Restoring the code would remove our injection, meaning our code can only run once. All current techniques (including `injectso`) use this method.

Since our main goal is to let our code run multiple times, we need a place in which our code can reside. We can store code on the stack and heap, but we can't guarantee execution. We can store code at `$eip`, but we are limited to a run-once scheme. We have an interesting dilemma on our hands: we have no place to put our code and guarantee repeated execution.

Spoiler alert: We will inject into a new anonymous memory mapping. This mapping will hold our malicious code and will be marked as executable. The next few chapters discuss how we go about forcing the victim process to create a new memory mapping.